

# CS 4530: Fundamentals of Software Engineering

## Module 11.1: Distributing Processing

---

Jon Bell, Adeel Bhutta, Mitch Wand  
Khoury College of Computer Sciences

# Learning Goals for this Lesson


---

- By the end of this lesson, you should be able to...
  - Recognize a few common software architectures
  - Discuss some of the tradeoffs of scalability, performance, and fault tolerance between these architectures

# Distributed Software Architectures

---

- Goal: abstract details away into reusable components
- Enables exploration of design alternatives
- Allows for analysis of high-level design before implementation
- Match system requirements to quality attributes of common architectural patterns



# Review: Challenges of Distributed Systems

---

- More machines mean more links that can fail
- Networks introduce delays
- Networks still fail, intermittently and for long periods
- Networks rely on fallible external administrators
- Sequential consistency is impossible

# Questions to Ask About Distributed Architectures?

---

- How many individual pieces can fail before the whole fails? Who is responsible for those pieces?
- How complicated is it...
  - To operate?
  - To debug?
  - To set up a development environment?
- How much CPU/RAM/bandwidth is needed to run it? (in total and per-node)
- What is the strategy for increasing capacity?

# A brief survey of distributed architectures


---

1. Monolithic server
2. Tiered architectures
3. Microservice architectures

# 1. The Monolith Architecture Relies on a Single Server

---


- Simplest answer to consistency problem: have only one server, one source of truth
- Still “distributed” in that we have many clients
- Sacrifices:
  - Scalability
  - Performance
  - Fault tolerance



# Monolithic Architectures Struggle to Scale

---

- Scalability - How to go from 10 to 100 to 1,000 clients?
- Performance - How to access 100's of GB of data concurrently?
- Fault tolerance - What if server crashes?






# Replication Alone is Not The Answer

---

- Constraints:
  - Latency: Speed of light ( $\sim 1\text{ns/ft}$ )
  - Throughput: Long-distance links between servers are relatively low throughput (10's of Gbps, compare to 100's of Gbps within a single server)
- Tradeoffs for replication, particularly over long distances:
  - Replication will *add* latency, not reduce it
  - Usually not enough bandwidth to maintain replication of all data across all nodes

## 2. Tiered Architectures

- Key idea: Partition the system into distinct tiers based on responsibilities
- Each tier scales independently of the others - .com need not know about .org
- Satisfying a single request may require multiple tiers
- DNS is a tiered architecture
  - Example: scale .com differently from .gov



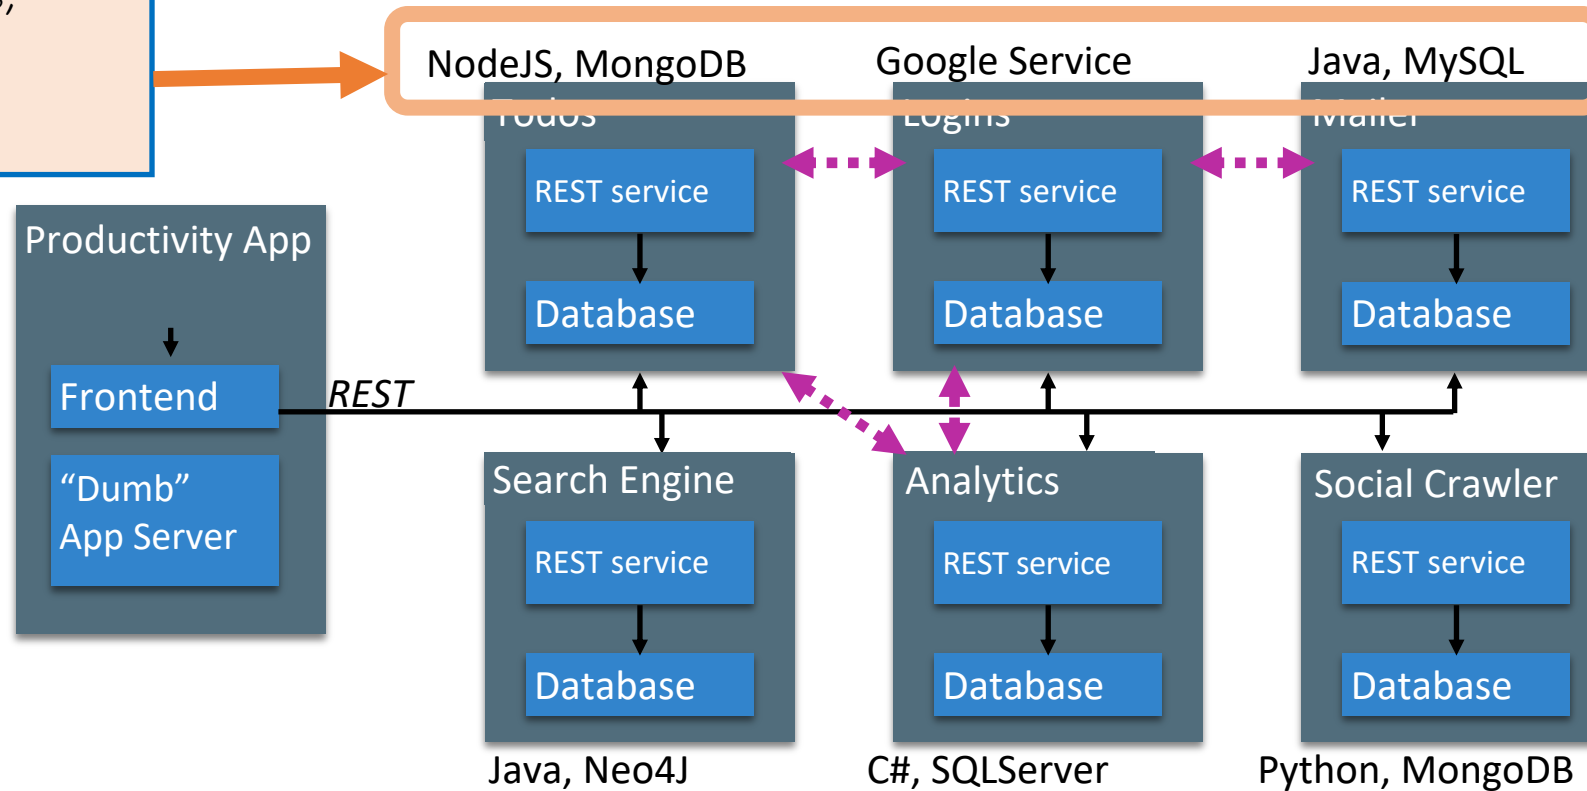
# 3. Microservice Architectures

---

- Organize implementation around components (responsibilities)
- Each component is implemented independently
- Each component is
  - independently replaceable,
  - independently updatable
- Components can be built as libraries, but more usually as web services
- Services communicate via a well-defined protocol (typically REST/http, though others are possible)

# Microservices: Schematic Example


Different languages,  
different operating  
systems



# Microservices are (a) highly scalable and (b) trendy

---

- Microservices at Netflix:
  - 100s of microservices
  - 1000s of daily production changes
  - 10,000s of instances
  - BUT:
  - only 10s of operations engineers




<https://medium.com/refraction-tech-everything/how-netflix-works-the-hugely-simplified-complex-stuff-that-happens-every-time-you-hit-play-3a40c9be254b>

# Microservice Advantages and Disadvantages

---

- Advantages
  - services may scale differently, so can be implemented on hardware and software appropriate for each
  - services are independent (yay for interfaces!) so can be developed and deployed independently
- Disadvantages
  - Shared data?
  - Requires high availability
  - Service discovery?
  - Data consistency?
  - Overall system complexity

# Microservices vs Monoliths



but remember the skill of the team will outweigh any monolith/microservice choice

<https://martinfowler.com/microservices/>

# Learning Goals for this Lesson

---

- You should now be able to
  - Recognize a few common software architectures
  - Discuss some of the tradeoffs of scalability, performance, and fault tolerance between these architectures